

Upcoming session:

July 11 – August 28

Enrolled


Computing regression parameters (gradient descent example)

The data

Consider the following 5 point synthetic data set:

1	X	Y
2	0	1
3	1	3
4	2	7
5	3	13
6	4	21

Which is plotted below:



What we need

Now that we've computed the regression line using a closed form solution let's do it again but with gradient descent.

Recall that:

- The derivative of the cost for the intercept is the sum of the errors
- The derivative of the cost for the slope is the sum of the product of the errors and the input

We will need a starting value for the slope and intercept, a `step_size` and a tolerance

- `initial_intercept = 0`
- `initial_slope = 0`
- `step_size = 0.05`
- `tolerance = 0.01`

The algorithm

In each step of the gradient descent we will do the following:

1. Compute the predicted values given the current slope and intercept
2. Compute the prediction errors (prediction - Y)
3. Update the intercept:
 - compute the derivative: `sum(errors)`
 - compute the adjustment as `step_size` times the derivative
 - decrease the intercept by the adjustment
4. Update the slope:
 - compute the derivative: `sum(errors*input)`
 - compute the adjustment as `step_size` times the derivative
 - decrease the slope by the adjustment
5. Compute the magnitude of the gradient
6. Check for convergence

The algorithm in action

First step:

Intercept = 0

Slope = 0

1. `predictions = [0, 0, 0, 0, 0]`
2. `errors = [-1, -3, -7, -13, -21]`
3. update Intercept
 - `sum([-1, -3, -7, -13, -21]) = -45`
 - `adjustment = 0.05 * 45 = -2.25`
 - `new_intercept = 0 - -2.25 = 2.25`

4. update Slope

- $\text{sum}([0, 1, 2, 3, 4] * [-1, -3, -7, -13, -21]) = -140$
- $\text{adjustment} = 0.05 * 45 = -7$
- $\text{new_slope} = 0 - -7 = 7$

5. $\text{magnitude} = \sqrt{(-45)^2 + (-140)^2} = 147.05$

6. $\text{magnitude} > \text{tolerance}$: not converged

Second step:

Intercept = 2.25

Slope = 7

1. $\text{predictions} = [2.25, 9.25, 16.25, 23.25, 30.25]$

2. $\text{errors} = [1.25, 6.35, 9.25, 10.25, 9.25]$

3. update Intercept

- $\text{sum}([1.25, 6.35, 9.25, 10.25, 9.25]) = 36.25$
- $\text{adjustment} = 0.05 * 36.25 = 1.8125$
- $\text{new_intercept} = 2.25 - 1.8125 = 0.4375$

4. update Slope

- $\text{sum}([0, 1, 2, 3, 4] * [1.25, 6.35, 9.25, 10.25, 9.25]) = 92.5$
- $\text{adjustment} = 0.05 * 92.5 = 4.625$
- $\text{new_slope} = 7 - 4.625 = 2.375$

5. $\text{magnitude} = \sqrt{(36.25)^2 + (92.5)^2} = 99.35$

6. $\text{magnitude} > \text{tolerance}$: not converged

Third step:

Intercept = 0.4375

Slope = 2.375

1. $\text{predictions} = [0.4375, 2.8125, 5.1875, 7.5625, 9.9375]$

2. $\text{errors} = [-0.5625, -0.1875, -1.8125, -5.4375, -11.0625]$

3. update Intercept

- $\text{sum}([-0.5625, -0.1875, -1.8125, -5.4375, -11.0625]) = -19.0625$
- $\text{adjustment} = 0.05 * -19.0625 = -0.953125$
- $\text{new_intercept} = 0.4375 - -0.953125 = 1.390625$

4. update Slope

- $\text{sum}([0, 1, 2, 3, 4] * [-0.5625, -0.1875, -1.8125, -5.4375, -11.0625]) = -64.375$
- $\text{adjustment} = 0.05 * -64.375 = -3.21875$
- $\text{new_slope} = 2.375 - 3.21875 = -0.84375$

5. $\text{magnitude} = \sqrt{(-19.0625)^2 + (-64.375)^2} = 67.13806$

6. $\text{magnitude} > \text{tolerance}$: not converged

Let's skip forward a few steps... after the 77th step we have gradient magnitude 0.0107.

78th Step:

Intercept = -0.9937

Slope = 4.9978

1. $\text{predictions} = [-0.99374, 4.00406, 9.00187, 13.99967, 18.99748]$

2. $\text{errors} = [-1.99374, 1.00406, 2.00187, 0.99967, -2.00252]$

3. update Intercept

- $\text{sum}([-1.99374, 1.00406, 2.00187, 0.99967, -2.00252]) = 0.009341224$
- $\text{adjustment} = 0.05 * 0.009341224 = 0.0004670612$
- $\text{new_intercept} = -0.9937 - 0.0004670612 = -0.994207$

4. update Slope

- $\text{sum}([0, 1, 2, 3, 4] * [-1.99374, 1.00406, 2.00187, 0.99967, -2.00252]) = -0.0032767$
- $\text{adjustment} = 0.05 * -0.0032767 = -0.00016383$
- $\text{new_slope} = 4.9978 - 0.00016383 = 4.9979$

5. $\text{magnitude} = \sqrt{()^2 + ()^2} = 0.0098992$

6. $\text{magnitude} < \text{tolerance}$: converged!


Final slope: -0.994

Final Intercept: 4.998


If you continue you will get to (-1, 5) but at this point the change in RSS (our cost) is negligible.

Visualizing the steps:


After the first step we have this line:




After the second step we have this line:



After the third step we have this line:



And after the final step we have this line:



Mark as completed

